

OCR Technology Scan

— 2026 —

*Landscape Review: VLMs • Managed APIs • Open-Source Engines • IDP
Platforms*

April 21, 2026

*Prepared from four parallel research streams
Snapshot date: April 20, 2026*

CONFIDENTIAL — Executive Distribution Only

Executive Summary

OCR as a distinct category has effectively dissolved into document understanding. In the last 12–18 months, vision-language models — both frontier (GPT-5, Claude 4.x, Gemini 3) and small document-specialized open-weight models (PaddleOCR-VL-1.5 at 0.9B, DeepSeek-OCR-2 at 3B, MinerU2.5 at 1.2B, olmOCR-2 at 7B) — have collapsed the classic layout-detect → recognize → reconstruct pipeline into a single forward pass that emits Markdown or JSON directly.

On public benchmarks like OmniDocBench v1.5/1.6, small open-weight document VLMs now match or exceed frontier APIs at 100× lower per-page cost. Pricing has compressed brutally on the managed side (Mistral OCR 3 at \$1–\$2 per 1,000 pages vs. \$30–\$65 for the hyperscalers’ form-parsing tiers), while the enterprise IDP category is being reshaped simultaneously by GenAI-first re-platformings and a consolidation wave.

For anyone choosing OCR tooling today, the right question is no longer “which OCR engine?” but “where on the stack do I want to sit?” — raw VLM in your own inference, managed API with structured JSON, self-hosted document pipeline, or full IDP workflow platform. The cost, accuracy, and control trade-offs between those four choices are now larger than the trade-offs inside any one of them.

Landscape Overview

Slice 1 — Foundation-Model / VLM-Based OCR

Frontier multimodal LLMs from OpenAI, Anthropic, Google, xAI, and Cohere can read documents end-to-end in a single pass; open-weight alternatives (Qwen3-VL, InternVL3, DeepSeek-OCR-2, MiniCPM-V 4.5, Llama 3.2 Vision, Molmo, Pixtral) have closed — and in some cases crossed — the gap on pure document benchmarks. The paradigm shift is that small, document-specialized VLMs post-trained on OCR/layout/table/formula objectives now beat generalist frontier models on OmniDocBench while running on a single consumer GPU. Hallucination remains the biggest reliability concern; Claude leads on hallucination rate (~0.09% on CC-OCR), GPT-5 on raw text edit distance, and Gemini 3 on price/performance. Cost curves have dropped roughly 80% since 2024.

Slice 2 — Cloud Hyperscaler & Managed OCR APIs

The “send bytes, receive structured text” category splits into two clear tiers. Hyperscalers (AWS Textract, Google Document AI, Azure AI Document Intelligence) compete on compliance, regional coverage, and prebuilt catalog breadth. API-first specialists (Mistral OCR 3, Reducto, LlamaParse, Unstructured.io, Upstage, Jina Reader) compete on price, Markdown/JSON fidelity, and raw accuracy. Mistral OCR 3 at \$1–\$2 per 1,000 pages and Unstructured Fast at \$1 per 1,000 pages make Textract AnalyzeDocument (\$65) look legacy. Agentic multi-pass OCR is a real product category for accuracy-critical workloads.

Slice 3 — Open-Source OCR Engines & Self-Hostable Libraries

Two distinct tiers. Classic CPU-friendly engines (Tesseract 5.5.2, PaddleOCR PP-OCRv5, EasyOCR, OCRmyPDF, Kraken, RapidOCR) remain valuable for clean text, edge devices, and air-gapped deployments but are no longer competitive on layout, tables, or formulas. The 2024–2026 wave of VLM-based document pipelines (Marker, Surya, olmOCR-2, MinerU, PaddleOCR-VL, Docling) dominates benchmarks — MinerU3 reports 95.69 on OmniDocBench v1.6. Licensing is now a first-class selection criterion: Docling (MIT) and MinerU (Apache-2.0 as of April 2026) are the commercial-safest high-quality choices.

Slice 4 — Intelligent Document Processing (IDP) Platforms

Gartner published its first IDP Magic Quadrant in September 2025, naming Hyperscience, ABBYY, Tungsten Automation, UiPath, and Infrd as Leaders. Two structural shifts dominate: (1) every platform has replaced rigid templates with prompt-tunable LLM/VLM extraction; and (2) a consolidation wave — UiPath acquired WorkFusion (Feb 2026), Sirion acquired Eigen Technologies, Klippa joined SER Group, AlphaSense acquired Tegos + Carousel. Vertical specialists remain strongly differentiated where domain data is specific.

Comparison Matrix — Leading Tools

Dimensions: Type, Deployment, Primary Strength, Notable Weakness, Cost (~2026), License / Lock-in, Best For. Tools ordered by slice.

Tool	Type	Deployment	Primary Strength	Notable Weakness	Cost (~2026)	License / Lock-in	Best For
GPT-5 / GPT-5.4	Frontier VLM	API (OpenAI, Azure)	Lowest edit distance on OmniDocBench among frontier APIs	Costly vs open-weight; no native bboxes	~\$1.75 / \$14 per 1M tokens	Proprietary, vendor lock-in	Hardest documents where cost is secondary
Claude Sonnet 4.6 / Opus 4.7	Frontier VLM	API (Anthropic, Bedrock, Vertex)	Lowest hallucination rate (~0.09%); 1M-token context; high-res vision	Higher image-token cost on long PDFs	Sonnet: \$3/\$15 per 1M tokens	Proprietary	Agentic doc workflows; compliance-sensitive
Gemini 3 Pro / 2.5 Flash	Frontier VLM	API (Google, Vertex)	Top OmniDocBench score; cheapest Flash tier	Grounding/bbox via prompting only	Flash among cheapest vision APIs	Proprietary	Long-document, high-volume, price-sensitive
Qwen3-VL / Qwen3.5-VL	Open-weight VLM	Self-host / hosted API	Open-weight SOTA; 32 languages; strong Asian scripts	235B variant heavy to self-host	Free; hosted ~\$0.10–\$0.30/1M tokens	Apache-2.0 / Qwen	Multilingual self-hosted document AI
DeepSeek-OCR-2 (3B)	Doc-specialized VLM	Self-host / hosted	SOTA small-model on OmniDocBench v1.5 (91.09)	Narrower general-VLM reasoning	Free; hosted ~\$0.15/1M tokens	Open weights	Dense layout / optical-compression
PaddleOCR-VL-1.5 (0.9B)	Doc-specialized VLM	Self-host	94.5% OmniDocBench v1.5; 109 languages; NaViT	Paddle ecosystem friction	Free	Apache-2.0	Multilingual, self-hosted, GPU-available
AWS Textract	Managed API	Managed (AWS)	Mature Forms/Tables/Queries; FedRAMP/HIPAA	Quality lags VLM-backed APIs	\$1.50–\$65 / 1k pages	Proprietary; AWS lock-in	Regulated US enterprise on AWS
Google Document AI	Managed API	Managed (GCP)	Gemini 3-backed Layout Parser; rich processor catalog	Processor sprawl; narrower data-residency	Layout \$10/1k; Form \$30/1k	Proprietary; GCP lock-in	Enterprises on Vertex AI / BigQuery
Azure AI Doc Intelligence	Managed API	Managed (Azure)	Prebuilt tax/mortgage/ID; Custom Generative Extraction	Dense SKU matrix	Read \$1.50; Prebuilt \$10; Custom \$30/1k	Proprietary; Azure lock-in	Microsoft/Azure shops; regulated
Mistral OCR 3	API specialist	Managed (multi-cloud)	88.9% handwriting, 96.6% tables; EU-hosted	Vendor-reported benchmarks	\$2/1k pages (\$1 Batch)	Proprietary, cross-cloud	High-accuracy at low cost; EU residency
Reducto	API (agentic)	Managed	Hybrid multi-pass OCR + VLM; 99.24% clinical SLAs	Slower than single-pass	\$0.015/page (~\$15/1k)	Proprietary	Enterprise accuracy-critical
LlamaParse v2	API specialist	Managed	Tiered modes; LlamaIndex-native; pinned versions	Credit-based pricing complexity	Fast ~\$1.25/1k; Agentic Plus ~\$12+/1k	Proprietary	RAG/agent pipelines on LlamaIndex
Unstructured.io (API)	API specialist	Managed	25+ file types; element-tree output; RAG-first	OCR quality = underlying engine	Fast \$1/1k; Hi-Res \$10/1k	Proprietary	RAG ingestion, heterogeneous corpora
Tesseract 5.5.2	OSS engine	Self-host (CPU)	100+ languages; runs anywhere; reference baseline	Weak on layout/tables/handwriting	Free	Apache-2.0	Clean printed text, offline/edge
PaddleOCR 3.4.1	OSS engine + VLM	Self-host	Most active classic OCR; PP-OCRv5 + VL unified	Paddle-framework-leaning	Free	Apache-2.0	Self-hosted multilingual OCR

Tool	Type	Deployment	Primary Strength	Notable Weakness	Cost (~2026)	License / Lock-in	Best For
Surya + Marker	OSS VLM pipeline	Self-host (GPU)	PDF→Markdown; tables/math/multi-column	GPL-3.0 requires Datalab license	Free (OSS) / paid (commercial)	GPL-3.0	Academic / technical PDF, non-commercial
olmOCR-2 (7B)	OSS VLM (Ai2)	Self-host (GPU)	Qwen2.5-VL fine-tune; 82.4 olmOCR-Bench	16–20 GB VRAM; English-leaning	Free	Apache-2.0	LLM pretraining corpus, scale OCR
MinerU 3.1	OSS VLM pipeline	Self-host (GPU)	95.69 on OmniDocBench v1.6; native DOCX	GPU-first; smaller English community	Free	Apache-based (Apr 2026)	Commercial self-hosted, top-quality
Docling 2.90	OSS pipeline (IBM)	Self-host (CPU+GPU)	MIT-licensed; pluggable VLMs; LF governance	Slightly below MinerU on benchmarks	Free	MIT	Commercial-safe enterprise RAG
Rossum Aurora	Enterprise IDP	SaaS / private cloud	Transactional T-LLM; template-free; ERP connectors	Narrow to AP/PO/orders	From ~\$1,500/mo	Proprietary	AP automation, transactional docs
ABBYY Vantage 3.0	Enterprise IDP	SaaS / on-prem	200+ skills; deep compliance; 35+ yr heritage	Heavy deployment; migration friction	Enterprise quote, mid-six-figure	Proprietary	Regulated enterprises, on-prem
Hyperscience Hypercell	Enterprise IDP	SaaS / on-prem / air-gapped	ORCA VLM; blocks-and-flows; gov-grade	Premium pricing; long implementation	\$250k–\$1M+ annual	Proprietary	Government, FSI, air-gapped HITL
UiPath Document Understanding	Enterprise IDP	SaaS	Tightest RPA binding; Autopilot for IXP	AI Unit pricing can escalate	1–2 AI Units/page	Proprietary	Shops already on UiPath RPA
Tungsten TotalAgility	Enterprise IDP	SaaS / on-prem	Kofax+Ephesoft rollup; wide ERP/RPA footprint	SKU confusion post-consolidation	\$100k+ annual floor	Proprietary	Large bank/gov, existing Kofax base
Nanonets	Mid-market IDP	SaaS	Low-code workflow builder; fast TTV	Depth gaps for regulated enterprise	\$0–\$999/mo	Proprietary	SMB/mid-market, rapid deployment
Docsumo	Mid-market IDP	SaaS	Lending/insurance focus; transparent pricing	Lighter integrations	Free (100 credits); Growth \$299/mo	Proprietary	Mid-market FS/insurance
Veryfi	Vertical (receipts)	SaaS / Mobile SDK	Sub-5s receipt extraction; mobile SDK	Narrow to receipts/invoices	Free tier; Enterprise from \$500/mo	Proprietary	Expense mgmt, field service
Mathpix	Vertical (STEM)	Managed API	Industry-standard math/chem OCR; LaTeX	Narrow to scientific docs	Per-snip / per-page quotes	Proprietary	Publishers, EdTech, STEM data
Ocrolus	Vertical (lending)	SaaS	Bank statements, tax returns, fraud detection	US lending-specific	Enterprise quote	Proprietary	Consumer/SMB lending
MyScript	Vertical (ink)	SDK / on-device	Stroke-based handwriting; 70+ languages	Not image OCR — digital ink only	OEM licensing	Proprietary	Note-taking apps, OEM

Deep Dive — Slice 1: Foundation-Model / VLM-Based OCR

The central finding: generalist frontier VLMs lead on breadth and ease-of-use, but small document-post-trained open-weight models now hold the top of every pure-OCR benchmark. Document-specific post-training beats scale.

Frontier Proprietary VLMs

- **OpenAI GPT-5 / GPT-5.4** — Lowest OmniDocBench edit distance (0.02) among frontier APIs. Handles dense scans, handwritten forms, engineering diagrams in a single pass. Pricing: \$1.75/\$14 per 1M tokens. Weakness: no native bboxes; occasional “silent correction” of typos.
- **Anthropic Claude Sonnet 4.6 / Opus 4.7** — The hallucination leader (~0.09% on CC-OCR). Opus 4.7 introduced high-resolution image support (2576 px). Sonnet 4.6 released Feb 2026; 1M-context beta since Aug 2025. Pricing: Sonnet \$3/\$15, Opus \$5/\$25. Strong for agentic document workflows.
- **Google Gemini 3 Pro / 2.5 Flash** — Gemini 3 Pro leads OmniDocBench v1.5 among frontier APIs (overall ED 0.115). Flash is the cheapest credible frontier vision API. Added as first-class OCR model in Tensorlake and Google Document AI Layout Parser.
- **xAI Grok 4 / Grok Vision** — Image analysis + OCR added July 2025. ~80–85% accuracy on straightforward docs; lags GPT-5/Claude on vision specifics. Narrower adoption.
- **Cohere Aya Vision / Command A Vision** — Multilingual across 23 languages. Aya Vision is open-weight but non-commercial (CC-BY-NC). Struggles on pure document-OCR prompts.

Open-Weight VLMs

- **Qwen3-VL / Qwen3.5-VL** — Open-weight SOTA; 32 languages; robust to blur/tilt. Qwen3.5 reportedly breaks 90 on OmniDocBench v1.5. Apache-2.0 / Qwen license.
- **InternVL 3 / 3.5** — InternVL3-78B hits OCRBench 906, DocVQA ~95. Supports 4K images. Fully open-weight.
- **DeepSeek-OCR-2 (3B)** — 91.09 on OmniDocBench v1.5 at 3B params; SOTA for small models. Runs on single consumer GPU.
- **Mistral Pixtral 12B** — Apache-2.0; strong ChartQA (83.7%); 128k context. Pairs with Mistral OCR 3 for production.
- **MiniCPM-V 4.5** — Edge-friendly; 46.7% of Qwen2.5-VL-7B's VRAM. Strong on-device OCR for English.
- **Meta Llama 3.2 Vision (11B/90B)** — Common OSS baseline (~82% OCR accuracy). Hallucinates on complex charts.
- **AI2 Molmo / PixMo** — Fully open (weights + data + code); 2D pointing grounding. CVPR 2025. Not document-specialized.

Slice 1 Takeaways

- Frontier API choice is primarily an error-profile choice: Claude (low hallucination) vs. GPT-5 (low edit distance) vs. Gemini (cheap and long).
- For self-hosters, PaddleOCR-VL-1.5, DeepSeek-OCR-2, and olmOCR-2 beat generalist frontier VLMs at a fraction of cost.
- Hallucination mitigation (MARINE at ICML 2025, contextual-embedding detectors) is the new accuracy frontier.
- OmniDocBench v1.5 is widely regarded as saturated; LlamaIndex's Real5-OmniDocBench and ParseBench are successors.

Deep Dive — Slice 2: Cloud Hyperscaler & Managed OCR APIs

Hyperscaler Document Services

- **AWS Textract** — DetectDocumentText, AnalyzeDocument (Forms/Tables/Queries/Signatures). June 2025 accuracy update. Pricing: \$1.50/1k (Detect) to \$65/1k (AnalyzeDocument). No built-in Markdown output.
- **Google Cloud Document AI** — Gemini 3-backed Layout Parser, 92–96% accuracy on complex tables, native DOCX/PPTX/XLSX support. Layout Parser \$10/1k pages. Broadest processor catalog.
- **Azure AI Document Intelligence** — Custom Generative Extraction (schema-in, JSON-out with citations) GA 2025. Read \$1.50/1k, Prebuilt \$10/1k, Custom \$30/1k.

API-First Specialists

- **Mistral OCR 3** — Purpose-built OCR VLM. Vendor-reported 88.9% handwriting, 96.6% tables. \$2/1k pages; \$1/1k Batch API. Multi-cloud distribution.
- **Reducto** — Agentic OCR: bounding-box detector + VLM verification. \$0.015/page. 99.24% on clinical SLAs. \$108M total funding. 1B+ pages processed.
- **LlamaParse v2** — Four tiers with swappable VLM backends. Credit pricing \$1.25/1k credits; Fast ≈ 1 credit/page, Agentic Plus 10+.
- **Unstructured.io** — 25+ file types; normalized element tree. Fast \$1/1k, Hi-Res \$10/1k. SOC 2 Type 2. Strong RAG preprocessing.
- **Upstage Document Parse** — Strong on Asian languages. Parse \$0.01/page; Extract +\$0.03/page. SOC 2; AWS Marketplace.
- **Jina Reader** — URL-to-Markdown; also handles PDFs. 10M free tokens per key. Best-fit for web-heavy RAG ingestion.

Slice 2 Takeaways

- VLMs are now the OCR backbone beneath both hyperscaler and specialist APIs.
- Mistral OCR 3 + Unstructured Fast redefine price-floor: \$1–2 per 1k pages.
- Structured Markdown/JSON + bbox + chunk IDs are the default output contract.
- Hyperscalers win on compliance; specialists win on price and hard-document accuracy.

Deep Dive — Slice 3: Open-Source OCR Engines & Libraries

Classic CPU-Friendly Engines

- **Tesseract 5.5.2** — Apache-2.0, ~73.6k stars. LSTM for 100+ languages. Maintenance-only cadence.
- **PaddleOCR 3.4.1** — Apache-2.0, ~76k stars. Unified PP-OCRv5 + PaddleOCR-VL-1.5 (0.9B VLM). 109 languages. 94.5% OmniDocBench.
- **EasyOCR 1.7.2** — Apache-2.0, ~29k stars. Simple Python API, 80+ languages. Last release ~19 months old.

- **OCRmyPDF 17.4.1** — MPL-2.0, ~33k stars. Canonical PDF→searchable-PDF pipeline. Actively maintained.
- **RapidOCR 3.8.1** — Apache-2.0. ONNX/OpenVINO wrapper. Python/C++/Java/C#/Go/JS bindings. Default embedded OCR for Docling.

VLM-Based Document Pipelines

- **Surya 0.17.1 + Marker 1.10.2** — GPL-3.0. PDF→Markdown/JSON. Monthly releases. Excellent on academic papers.
- **olmOCR-2 (7B)** — Apache-2.0. Qwen2.5-VL-7B fine-tune. 82.4 on olmOCR-Bench, beating Marker (76.1) and MinerU (75.8). GRPO RL training.
- **MinerU 3.1.0** — Apache-based (relicensed Apr 2026 from AGPL-3.0). ~60k stars. 95.69 on OmniDocBench v1.6. Extremely active.
- **Docling 2.90** — MIT, ~58k stars. IBM Research / LF governance. Pluggable OCR + table models. LangChain/LlamaIndex integrations.
- **Unstructured OSS 0.22.21** — Apache-2.0, ~14.5k stars. Broad “any doc → chunked elements” ETL for RAG.

Slice 3 Takeaways

- VLM-based pipelines top OmniDocBench and beat frontier APIs at 100× lower cost for self-hosters.
- Licensing is first-class: **Docling (MIT)**, **MinerU (Apache)**, **Tesseract/PaddleOCR/olmOCR (Apache-2.0)** are commercial-safe.
- GPU-first vs CPU-first has split cleanly: classic engines serve edge/offline; VLM pipelines serve server workloads.

Deep Dive — Slice 4: IDP Platforms & Vertical Specialists

Enterprise IDP (Gartner 2025 Leaders)

- **Hyperscience** — Blocks-and-flows architecture; ORCA VLM; government/FSI. \$250k–\$1M+ annually. “Intelligent Inference” Spring 2026 rebrand.
- **ABBY Vantage 3.0** — 35+ year heritage, 200+ skills. Direct Azure OpenAI integration. Strong on-prem story.
- **Tungsten TotalAgility** — Full-stack capture + RPA + orchestration. Platform unified launch Sept 2025. GenAI Copilots.
- **UiPath Document Understanding** — Autopilot for IXP; Generative Validation. Feb 2026: acquired WorkFusion for financial-crime vertical.
- **Infrrd** — Inaugural 2025 Gartner Leader; mortgage/insurance strength; HITL guarantees.
- **Rossum Aurora** — Transactional-specific LLM (DocILE-trained); template-free; direct ERP connectors. 500+ enterprise customers.

Mid-Market IDP & Vertical Specialists

- **Nanonets** — Low-code workflow builder; consumption pricing. \$0–\$999/mo.
- **Docsumo** — Lending/insurance focus. Free 100 credits, Growth \$299/mo.
- **Veryfi** — Receipts/invoices; sub-5s mobile SDK + API; fraud detection.

- **Mathpix** — STEM/math/chem OCR; LaTeX output; SuperNet-108 deployed 2025.
- **Ocrolus** — Consumer/SMB lending; bank statements, tax returns, fraud. 175+ lenders. Encore platform (Oct 2025).
- **MyScript** — Digital ink SDK; 70+ languages; on-device. SDK 4.3 (Jan 2026).

Slice 4 Takeaways

- LLM/VLM extraction is replacing hand-built templates. Schema-from-samples + prompt-tunable extraction is the new floor.
- Consolidation wave: UiPath+WorkFusion, Sirion+Eigen, SER+Klippa, Tungsten+Kofax+Ephesoft, AlphaSense+Tegus+Carousel.
- IDP moat is shifting from extraction accuracy (commoditizing) to workflow, HITL, governance, and ERP integration depth.
- Identity verification is a separate fraud-ML arms race against deepfakes — not an OCR problem.

Recommendations by Use Case

1. RAG pipeline over messy internal PDFs (commercial, self-hosted)

Start with **Docling** (MIT) for breadth or **MinerU 3.1** (Apache) for quality. Add **RapidOCR** as fallback. Avoid Marker/Surya unless you license Datalab commercially.

2. RAG pipeline fast, willing to pay per page

LlamaParse v2 (Cost-Effective) or **Unstructured.io Fast**. For harder docs, **Mistral OCR 3** at \$2/1k pages. **Reducto** if 99%+ accuracy SLAs needed.

3. Extract invoices/POs/receipts into ERP at scale

Rossum Aurora (T-LLM, ERP connectors) or **SAP Document AI** on S/4HANA. **Veryfi** for mobile. **Nanonets/Docsumo** for mid-market.

4. Government / regulated / air-gapped

Hyperscience (ORCA VLM, air-gapped) or **ABBYY Vantage 3.0** (on-prem DNA). **Azure AI Doc Intelligence** for Azure Government.

5. STEM / academic papers → LaTeX

Mathpix (managed, industry-standard) or **Marker + olmOCR-2** (self-hosted, GPL caveat on Marker).

6. Historical manuscripts / RTL / medieval scripts

Kraken (Apache-2.0) in the eScriptorium ecosystem.

7. Mobile app OCR without cloud dependency

MyScript (digital ink), **MiniCPM-V 4.5** (on-device VLM), or **RapidOCR** (CoreML/ONNX).

8. Maximize accuracy, cost is secondary

Reducto (agentic multi-pass), **Azure Custom Generative Extraction**, or **LlamaParse Agentic Plus**. Raw-VLM: **GPT-5.4** for edit distance, **Claude Opus 4.7** for low hallucination.

9. Minimize cost, self-host on a single GPU

DeepSeek-OCR-2 (3B) or **PaddleOCR-VL-1.5 (0.9B)** — both run on a single consumer GPU.

10. Identity verification for KYC/onboarding

Entrust/Onfido, Jumio, or Sumsu. Separate category — pick on fraud-ML maturity.

11. Replace aging Tesseract pipeline

RapidOCR (CPU, cross-language bindings). **PaddleOCR PP-OCRv5** for GPU. **Docling** to add table/formula handling.

Sources and References

This report synthesizes four parallel research streams run on 2026-04-20. Pricing reflects published vendor rate cards. Benchmark figures are cited from source; several are vendor-reported and should be verified before procurement decisions.

Slice 1 — VLM-based OCR

- CodeSOTA — Best OCR Models 2026 & OmniDocBench Leaderboard
- OpenDataLab OmniDocBench (CVPR 2025)
- LlamaIndex — “OmniDocBench is Saturated, What’s Next”
- OpenAI Cookbook — GPT-5.4 Multimodal Document Tips
- Claude API Pricing & Vision Docs (platform.claude.com)
- Qwen3-VL, InternVL3, DeepSeek-OCR-2, MiniCPM-V 4.5, Molmo — GitHub/arXiv
- ICML 2025 — MARINE Hallucination Mitigation

Slice 2 — Managed APIs

- AWS Textract, Google Document AI, Azure AI Document Intelligence — Official pricing
- Mistral OCR 3 announcement & pricing (mistral.ai)
- Reducto — Series B funding, hybrid architecture deep dive
- LlamaParse v2, Unstructured.io, Upstage, Jina Reader — Official docs

Slice 3 — Open-Source

- Tesseract, PaddleOCR, EasyOCR, OCRmyPDF, Kraken, RapidOCR, docTR — GitHub repos
- Surya, Marker, olmOCR-2, MinerU 3.1, Docling 2.90 — GitHub repos & announcements

Slice 4 — IDP & Verticals

- Gartner Magic Quadrant for IDP Solutions (Sept 2025)
- ABBYY, Hyperscience, UiPath, Tungsten — Press releases
- Rossum Aurora, Automation Anywhere, SAP Business AI — Official announcements
- Acquisition coverage: UiPath+WorkFusion, Sirion+Eigen, AlphaSense+Tegus
- WEF 2026 — Digital Identity + Deepfakes report